

Education

Bennett University

Bachelor of Technology in Computer Science and Engineering (Data Science) | CGPA: **8.5**

Greater Noida

2023-2027

Experience

Techknowledgehub

(Gurgaon)

Product Engineering Intern

Jun 2026 – Present

- Driving end-to-end product development across web and mobile, designing features and user workflows while building with Next.js/React, Expo/React Native, FastAPI, PostgreSQL, Redis, and Docker
- Diagnosing issues across services during development by tracing logs, error messages, and API responses to isolate root causes ahead of production rollout

Research Work

(Bennett University)

IST Computing

Feb 2026 – Present

- Simulating a multi-layer adaptive IDS across IoT → Cloud in OMNeT++, evaluating latency, throughput, and packet loss across **4 tiers** to identify optimal detection thresholds

GitHub Repository

Abhiyan

(Remote)

Data Analytics Intern

Oct 2025 – Dec 2025

- Analyzed sales data across **6 product categories** using SQL and Pandas, surfacing revenue trends and category-level performance for monthly stakeholder reporting
- Cleaned and transformed raw datasets, reducing null records by **~30%** and improving downstream reporting accuracy

Valentino Tech. Pvt. Ltd.

(Remote)

Business Development Intern

Jan 2025 – Feb 2025

- Built **3 executive-facing Power BI and Tableau dashboards** monitoring revenue, lead conversion, customer retention, and other key business KPIs for data-driven decision making
- Cleaned, transformed, and integrated data from **5 disparate sources**, reducing reporting discrepancies by **~25%** and streamlining weekly reporting workflows

Projects

DocuMind — RAG & Agentic AI System | *Python, FastAPI, PostgreSQL, Redis, Celery, Pinecone, LLMs*

GitHub

- Built a multi-layer AI system: RAG retrieval (chunking, embeddings, vector search via Pinecone), a SQL tool-calling layer combining unstructured documents with structured databases, and an agentic layer for dynamic tool selection and multi-step reasoning
- Instrumented evaluation metrics, observability, and monitoring at each layer, tracking request routing, latency, and RAGAS-based quality to catch degradation early
- Added guardrails and semantic caching for reliability, with async job queues (Celery + Redis) and REST APIs over a PostgreSQL backend

Real-Time Data Streaming Pipeline | *Kafka, Flink, AWS (EC2, S3, Glue, Athena)*

GitHub

- Designed Kafka-based streaming pipeline with 3 topics and consumer groups, simulating **~5,000 events/sec**
- Implemented Flink processing achieving **sub-second latency** for real-time transformations and aggregations
- Stored processed data in S3 (**~2GB/day**) and enabled querying via Athena with **<3s** response time

Financial Market Analytics Pipeline | *Airflow, dbt, PostgreSQL, Docker*

GitHub

- Built batch ETL pipeline ingesting OHLCV data for **10 stocks (~50K rows/month)** into PostgreSQL
- Designed dbt models (staging → marts) computing moving averages, returns, and volatility using SQL
- Containerized the pipeline using Docker Compose, reducing setup time from **~2 hours to <5 minutes**

Skills

Languages: Python, C++, SQL

Web & Backend: Next.js/React, React Native (Expo), FastAPI, Flask, REST APIs, Redis, Celery

Data & Analytics: Pandas, PostgreSQL, MySQL, dbt, Apache Airflow, ETL, Power BI, Tableau

AI & Streaming: RAG Pipelines, LLM Integration, Pinecone, Apache Kafka, Apache Flink

Cloud & Tools: AWS (EC2, S3, Glue, Athena), Docker, Git, Linux